

```

hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin$ hadoop jar square.jar squareNumbers.SQDriver
/kms/input.txt /kms/output/
22/07/11 15:22:56 INFO Configuration.deprecation: session.id is deprecated. Instead, use
dfs.metrics.session-id
22/07/11 15:22:56 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
22/07/11 15:22:56 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker,
sessionId= - already initialized
Exception in thread "main" org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory
hdfs://localhost:54310/kms/output already exists
    at org.apache.hadoop.mapred.FileOutputFormat.checkOutputSpecs(FileOutputFormat.java:132)
    at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:564)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:432)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1296)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1293)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1293)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:562)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:557)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1628)
    at org.apache.hadoop.mapred.JobClient.submitJobInternal(JobClient.java:557)
    at org.apache.hadoop.mapred.JobClient.submitJob(JobClient.java:548)
    at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:833)
    at squareNumbers.SQDriver.run(SQDriver.java:34)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
    at squareNumbers.SQDriver.main(SQDriver.java:41)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin$ hadoop jar square.jar squareNumbers.SQDriver
/kms/input.txt /kms/output2/
22/07/11 15:23:18 INFO Configuration.deprecation: session.id is deprecated. Instead, use
dfs.metrics.session-id
22/07/11 15:23:18 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
22/07/11 15:23:18 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with processName=JobTracker,
sessionId= - already initialized
22/07/11 15:23:18 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed.
Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/07/11 15:23:18 INFO mapred.FileInputFormat: Total input paths to process : 1
22/07/11 15:23:18 INFO mapreduce.JobSubmitter: number of splits:1
22/07/11 15:23:18 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local319021142_0001
22/07/11 15:23:18 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
22/07/11 15:23:18 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/07/11 15:23:18 INFO mapreduce.Job: Running job: job_local319021142_0001

```

```
22/07/11 15:23:18 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapred.FileOutputCommitter
22/07/11 15:23:18 INFO mapred.LocalJobRunner: Waiting for map tasks
22/07/11 15:23:18 INFO mapred.LocalJobRunner: Starting task: attempt_local319021142_0001_m_000000_0
22/07/11 15:23:18 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/07/11 15:23:18 INFO mapred.MapTask: Processing split: hdfs://localhost:54310/kms/input.txt:0+14
22/07/11 15:23:18 INFO mapred.MapTask: numReduceTasks: 1
22/07/11 15:23:18 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
22/07/11 15:23:18 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
22/07/11 15:23:18 INFO mapred.MapTask: soft limit at 83886080
22/07/11 15:23:18 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
22/07/11 15:23:18 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
22/07/11 15:23:18 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/07/11 15:23:18 INFO mapred.LocalJobRunner:
22/07/11 15:23:18 INFO mapred.MapTask: Starting flush of map output
22/07/11 15:23:18 INFO mapred.MapTask: Spilling map output
22/07/11 15:23:18 INFO mapred.MapTask: bufstart = 0; bufend = 56; bufvoid = 104857600
22/07/11 15:23:18 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214372(104857488);
length = 25/6553600
22/07/11 15:23:18 INFO mapred.MapTask: Finished spill 0
22/07/11 15:23:18 INFO mapred.Task: Task:attempt_local319021142_0001_m_000000_0 is done. And is in
the process of committing
22/07/11 15:23:18 INFO mapred.LocalJobRunner: hdfs://localhost:54310/kms/input.txt:0+14
22/07/11 15:23:18 INFO mapred.Task: Task 'attempt_local319021142_0001_m_000000_0' done.
22/07/11 15:23:18 INFO mapred.LocalJobRunner: Finishing task: attempt_local319021142_0001_m_000000_0
22/07/11 15:23:18 INFO mapred.LocalJobRunner: map task executor complete.
22/07/11 15:23:18 INFO mapred.LocalJobRunner: Waiting for reduce tasks
22/07/11 15:23:18 INFO mapred.LocalJobRunner: Starting task: attempt_local319021142_0001_r_000000_0
22/07/11 15:23:18 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/07/11 15:23:18 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin:
org.apache.hadoop.mapreduce.task.reduce.Shuffle@559edab0
22/07/11 15:23:18 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=334338464,
maxSingleShuffleLimit=83584616, mergeThreshold=220663392, ioSortFactor=10,
memToMemMergeOutputsThreshold=10
22/07/11 15:23:18 INFO reduce.EventFetcher: attempt_local319021142_0001_r_000000_0 Thread started:
EventFetcher for fetching Map Completion Events
22/07/11 15:23:18 INFO reduce.LocalFetcher: localfetcher#1 about to shuffle output of map
attempt_local319021142_0001_m_000000_0 decomp: 72 len: 76 to MEMORY
22/07/11 15:23:18 INFO reduce.InMemoryMapOutput: Read 72 bytes from map-output for
attempt_local319021142_0001_m_000000_0
22/07/11 15:23:18 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 72,
inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory ->72
22/07/11 15:23:18 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
22/07/11 15:23:18 INFO mapred.LocalJobRunner: 1 / 1 copied.
22/07/11 15:23:18 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0
on-disk map-outputs
22/07/11 15:23:18 INFO mapred.Merger: Merging 1 sorted segments
22/07/11 15:23:18 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total
size: 66 bytes
22/07/11 15:23:18 INFO reduce.MergeManagerImpl: Merged 1 segments, 72 bytes to disk to satisfy
reduce memory limit
22/07/11 15:23:18 INFO reduce.MergeManagerImpl: Merging 1 files, 76 bytes from disk
```

```
22/07/11 15:23:18 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
22/07/11 15:23:18 INFO mapred.Merger: Merging 1 sorted segments
22/07/11 15:23:18 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total
size: 66 bytes
22/07/11 15:23:18 INFO mapred.LocalJobRunner: 1 / 1 copied.
22/07/11 15:23:18 INFO mapred.Task: Task:attempt_local319021142_0001_r_000000_0 is done. And is in
the process of committing
22/07/11 15:23:18 INFO mapred.LocalJobRunner: 1 / 1 copied.
22/07/11 15:23:18 INFO mapred.Task: Task attempt_local319021142_0001_r_000000_0 is allowed to commit
now
22/07/11 15:23:18 INFO output.FileOutputCommitter: Saved output of task
'attempt_local319021142_0001_r_000000_0' to
hdfs://localhost:54310/kms/output2/_temporary/0/task_local319021142_0001_r_000000
22/07/11 15:23:18 INFO mapred.LocalJobRunner: reduce > reduce
22/07/11 15:23:18 INFO mapred.Task: Task 'attempt_local319021142_0001_r_000000_0' done.
22/07/11 15:23:18 INFO mapred.LocalJobRunner: Finishing task: attempt_local319021142_0001_r_000000_0
22/07/11 15:23:18 INFO mapred.LocalJobRunner: reduce task executor complete.
22/07/11 15:23:19 INFO mapreduce.Job: Job job_local319021142_0001 running in uber mode : false
22/07/11 15:23:19 INFO mapreduce.Job: map 100% reduce 100%
22/07/11 15:23:19 INFO mapreduce.Job: Job job_local319021142_0001 completed successfully
22/07/11 15:23:19 INFO mapreduce.Job: Counters: 38
File System Counters
FILE: Number of bytes read=8744
FILE: Number of bytes written=507910
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=28
HDFS: Number of bytes written=8
HDFS: Number of read operations=13
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
Map-Reduce Framework
Map input records=1
Map output records=7
Map output bytes=56
Map output materialized bytes=76
Input split bytes=88
Combine input records=0
Combine output records=0
Reduce input groups=1
Reduce shuffle bytes=76
Reduce input records=7
Reduce output records=1
Spilled Records=14
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=1
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=472907776
```

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Input Format Counters

Bytes Read=14

File Output Format Counters

Bytes Written=8

0

```
hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin$ hdfs dfs -ls /kms
```

Found 2 items

```
-rw-r--r--  1 hduser supergroup      14 2022-07-11 15:14 /kms/input.txt
```

```
drwxr-xr-x  - hduser supergroup      0 2022-07-11 15:19 /kms/output2
```

```
hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin$ hdfs dfs -ls /kms/output2
```

Found 2 items

```
-rw-r--r--  1 hduser supergroup      0 2022-07-11 15:19 /kms/output2/_SUCCESS
```

```
-rw-r--r--  1 hduser supergroup    28 2022-07-11 15:19 /kms/output2/part-00000
```

```
hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin$ hdfs dfs -cat /kms/output2/part-00000
```

```
sum      140
```

```
hduser@bmsce-Precision-T1700:/usr/local/hadoop/sbin$
```