

```

hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -mkdir /khushil_topn
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -put ./input.txt /khushil_topn/
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hdfs dfs -ls /khushil_topn/
Found 1 items
-rw-r--r--  1 hduser supergroup      103 2022-06-27 15:43 /khushil_topn/input.txt
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hadoop jar topn.jar TopNDriver
/khushil_topn/input.txt /khushil_topn/output
Exception in thread "main" java.lang.ClassNotFoundException: TopNDriver
    at java.net.URLClassLoader.findClass(URLClassLoader.java:382)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:418)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:351)
    at java.lang.Class.forName0(Native Method)
    at java.lang.Class.forName(Class.java:348)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:214)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
hduser@bmsce-Precision-T1700:~/Desktop/temperature$ hadoop jar topn.jar topn.TopNDriver
/khushil_topn/input.txt /khushil_topn/output
22/06/27 15:45:22 INFO Configuration.deprecation: session.id is deprecated. Instead, use
dfs.metrics.session-id
22/06/27 15:45:22 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker,
sessionId=
22/06/27 15:45:22 INFO input.FileInputFormat: Total input paths to process : 1
22/06/27 15:45:22 INFO mapreduce.JobSubmitter: number of splits:1
22/06/27 15:45:22 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local691635730_0001
22/06/27 15:45:22 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
22/06/27 15:45:22 INFO mapreduce.Job: Running job: job_local691635730_0001
22/06/27 15:45:22 INFO mapred.LocalJobRunner: OutputCommitter set in config null
22/06/27 15:45:22 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Waiting for map tasks
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Starting task: attempt_local691635730_0001_m_000000_0
22/06/27 15:45:22 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
22/06/27 15:45:22 INFO mapred.MapTask: Processing split:
hdfs://localhost:54310/khushil_topn/input.txt:0+103
22/06/27 15:45:22 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
22/06/27 15:45:22 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
22/06/27 15:45:22 INFO mapred.MapTask: soft limit at 83886080
22/06/27 15:45:22 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
22/06/27 15:45:22 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
22/06/27 15:45:22 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask$MapOutputBuffer
22/06/27 15:45:22 INFO mapred.LocalJobRunner:
22/06/27 15:45:22 INFO mapred.MapTask: Starting flush of map output
22/06/27 15:45:22 INFO mapred.MapTask: Spilling map output
22/06/27 15:45:22 INFO mapred.MapTask: bufstart = 0; bufend = 187; bufvoid = 104857600
22/06/27 15:45:22 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214316(104857264);
length = 81/6553600
22/06/27 15:45:22 INFO mapred.MapTask: Finished spill 0
22/06/27 15:45:22 INFO mapred.Task: Task:attempt_local691635730_0001_m_000000_0 is done. And is in
the process of committing
22/06/27 15:45:22 INFO mapred.LocalJobRunner: map
22/06/27 15:45:22 INFO mapred.Task: Task 'attempt_local691635730_0001_m_000000_0' done.
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Finishing task: attempt_local691635730_0001_m_000000_0
22/06/27 15:45:22 INFO mapred.LocalJobRunner: map task executor complete.
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Waiting for reduce tasks
22/06/27 15:45:22 INFO mapred.LocalJobRunner: Starting task: attempt_local691635730_0001_r_000000_0
22/06/27 15:45:22 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]

```

22/06/27 15:45:22 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin:  
org.apache.hadoop.mapreduce.task.reduce.Shuffle@40a5e65a  
22/06/27 15:45:22 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=334338464,  
maxSingleShuffleLimit=83584616, mergeThreshold=220663392, ioSortFactor=10,  
memToMemMergeOutputsThreshold=10  
22/06/27 15:45:22 INFO reduce.EventFetcher: attempt\_local691635730\_0001\_r\_0000000\_0 Thread started:  
EventFetcher for fetching Map Completion Events  
22/06/27 15:45:22 INFO reduce.LocalFetcher: localfetcher#1 about to shuffle output of map  
attempt\_local691635730\_0001\_m\_0000000\_0 decomp: 231 len: 235 to MEMORY  
22/06/27 15:45:22 INFO reduce.InMemoryMapOutput: Read 231 bytes from map-output for  
attempt\_local691635730\_0001\_m\_0000000\_0  
22/06/27 15:45:22 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 231,  
inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory ->231  
22/06/27 15:45:22 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning  
22/06/27 15:45:22 INFO mapred.LocalJobRunner: 1 / 1 copied.  
22/06/27 15:45:22 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0  
on-disk map-outputs  
22/06/27 15:45:22 INFO mapred.Merger: Merging 1 sorted segments  
22/06/27 15:45:22 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total  
size: 226 bytes  
22/06/27 15:45:22 INFO reduce.MergeManagerImpl: Merged 1 segments, 231 bytes to disk to satisfy  
reduce memory limit  
22/06/27 15:45:22 INFO reduce.MergeManagerImpl: Merging 1 files, 235 bytes from disk  
22/06/27 15:45:22 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce  
22/06/27 15:45:22 INFO mapred.Merger: Merging 1 sorted segments  
22/06/27 15:45:22 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total  
size: 226 bytes  
22/06/27 15:45:22 INFO mapred.LocalJobRunner: 1 / 1 copied.  
22/06/27 15:45:22 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use  
mapreduce.job.skiprecords  
22/06/27 15:45:23 INFO mapred.Task: Task:attempt\_local691635730\_0001\_r\_0000000\_0 is done. And is in  
the process of committing  
22/06/27 15:45:23 INFO mapred.LocalJobRunner: 1 / 1 copied.  
22/06/27 15:45:23 INFO mapred.Task: Task attempt\_local691635730\_0001\_r\_0000000\_0 is allowed to commit  
now  
22/06/27 15:45:23 INFO output.FileOutputCommitter: Saved output of task  
'attempt\_local691635730\_0001\_r\_0000000\_0' to  
hdfs://localhost:54310/khushil\_topn/output/\_temporary/0/task\_local691635730\_0001\_r\_0000000  
22/06/27 15:45:23 INFO mapred.LocalJobRunner: reduce > reduce  
22/06/27 15:45:23 INFO mapred.Task: Task 'attempt\_local691635730\_0001\_r\_0000000\_0' done.  
22/06/27 15:45:23 INFO mapred.LocalJobRunner: Finishing task: attempt\_local691635730\_0001\_r\_0000000\_0  
22/06/27 15:45:23 INFO mapred.LocalJobRunner: reduce task executor complete.  
22/06/27 15:45:23 INFO mapreduce.Job: Job job\_local691635730\_0001 running in uber mode : false  
22/06/27 15:45:23 INFO mapreduce.Job: map 100% reduce 100%  
22/06/27 15:45:23 INFO mapreduce.Job: Job job\_local691635730\_0001 completed successfully  
22/06/27 15:45:23 INFO mapreduce.Job: Counters: 38

#### File System Counters

FILE: Number of bytes read=18078  
FILE: Number of bytes written=516697  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=206  
HDFS: Number of bytes written=105  
HDFS: Number of read operations=13  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=4

#### Map-Reduce Framework

Map input records=6  
Map output records=21  
Map output bytes=187  
Map output materialized bytes=235  
Input split bytes=110  
Combine input records=0  
Combine output records=0  
Reduce input groups=15  
Reduce shuffle bytes=235  
Reduce input records=21  
Reduce output records=15  
Spilled Records=42  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=42  
CPU time spent (ms)=0  
Physical memory (bytes) snapshot=0  
Virtual memory (bytes) snapshot=0  
Total committed heap usage (bytes)=578289664  
Shuffle Errors  
BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0  
File Input Format Counters  
Bytes Read=103  
File Output Format Counters  
Bytes Written=105

hduser@bmsce-Precision-T1700:~/Desktop/temperature\$ hdfs dfs -ls /khushil\_topn/output/

Found 2 items

-rw-r--r--	1	hduser	supergroup	0	2022-06-27 15:45	/khushil_topn/output/_SUCCESS
-rw-r--r--	1	hduser	supergroup	105	2022-06-27 15:45	/khushil_topn/output/part-r-00000

hduser@bmsce-Precision-T1700:~/Desktop/temperature\$ hdfs dfs -cat /khushil\_topn/output/part-r-00000

hadoop 4

i3

am 2

hi 1

im 1

is 1

there 1

bye 1

learing 1

awesome 1

love 1

khushil 1

cool 1

and 1

using 1

hduser@bmsce-Precision-T1700:~/Desktop/temperature\$