

```
khushil@khushil-VirtualBox:~/Desktop/cassandra/bin$ sudo apt install cassandra
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  chromium-codecs-ffmpeg-extra gstreamer1.0-vaapi
  libgstreamer-plugins-bad1.0-0 libva-wayland2
Use 'sudo apt autoremove' to remove them.
Suggested packages:
  cassandra-tools
The following NEW packages will be installed:
  cassandra
0 upgraded, 1 newly installed, 0 to remove and 312 not upgraded.
Need to get 47.5 MB of archives.
After this operation, 58.8 MB of additional disk space will be used.
Get:1 https://downloads.apache.org/cassandra/debian 40x/main amd64 cassandra all 4.0.4 [47.5 MB]
Fetched 47.5 MB in 8s (5,919 kB/s)
Selecting previously unselected package cassandra.
(Reading database ... 191150 files and directories currently installed.)
Preparing to unpack .../cassandra_4.0.4_all.deb ...
Unpacking cassandra (4.0.4) ...
Setting up cassandra (4.0.4) ...
Adding group `cassandra' (GID 134) ...
Done.
vm.max_map_count = 1048575
net.ipv4.tcp_keepalive_time = 300
update-rc.d: warning: start and stop actions are no longer supported; falling ba
ck to defaults
Processing triggers for systemd (245.4-4ubuntu3.15) ...
khushil@khushil-VirtualBox:~/Desktop/cassandra/bin$ sudo systemctl enable cassandra
cassandra.service is not a native service, redirecting to systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable cassandra
khushil@khushil-VirtualBox:~/Desktop/cassandra/bin$ sudo systemctl start cassandra
khushil@khushil-VirtualBox:~/Desktop/cassandra/bin$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.0.0 | Cassandra 3.11.13 | CQL spec 3.4.4 | Native protocol v4]
cqlsh> create keyspace Employee with replication = {'class': 'SimpleStrategy', 'replication_factor': 2};
```

Warnings :

Your replication factor 2 for keyspace employee is higher than the number of nodes 1

```
cqlsh> create table EmployeeInfo (
... empid int,
... empname text,
... designation text,
... date_of_joining timestamp,
... salary double,
... dept text,
... );
```

**InvalidRequest: Error from server: code=2200 [Invalid query] message="No keyspace has been specified. USE a keyspace, or explicitly specify keyspace.tablename"**

```
cqlsh> use Employee;
```

```
cqlsh:employee> create table EmployeeInfo ( empid int, empname text, designation text, date_of_joining timestamp, salary double, dept text,
);
```

**InvalidRequest: Error from server: code=2200 [Invalid query] message="No PRIMARY KEY specified (exactly one required)"**

```
cqlsh:employee> create table EmployeeInfo ( empid int, empname text, designation text, date_of_joining timestamp, salary double, dept text, );
```

**InvalidRequest: Error from server: code=2200 [Invalid query] message="No PRIMARY KEY specified (exactly one required)"**

```
cqlsh:employee> create table EmployeeInfo ( empid int, empname text, designation text, date_of_joining timestamp, salary double, dept text, )
```

...

```
cqlsh:employee>
```

```
cqlsh:employee> create table EmployeeInfo ( empid int, empname text, designation text, date_of_joining timestamp, salary double, dept text, ... primary key(empid));
```

```
cqlsh:employee> begin batch
```

```
... insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (1, "ramesh", "Software Engineer II", "2022-05-05", 15000.0, "IT")
```

```
... insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (2, "suresh", "Software Engineer I", "2022-05-06", 10000.0, "IT")
```

```
... insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (3, "raju", "HR", "2022-06-06", 20000.0, "HR")
```

```
... apply batch;
```

**SyntaxException: line 2:105 no viable alternative at input ';' (... dept) values (1, ["ramesh"],...)**

```
cqlsh:employee> begin batch insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (1, 'ramesh', 'Software Engineer II', '2022-05-05', 15000.0, 'IT') insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (2, 'suresh', 'Software Engineer I', '2022-05-06', 10000.0, 'IT') insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (3, 'raju', 'HR', '2022-06-06', 20000.0, 'HR') apply batch;
```

```
cqlsh:employee> select * from EmployeeInfo
```

... ;

empid	date_of_joining	dept	designation	empname	salary
1	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	ramesh	15000
2	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	suresh	10000
3	2022-06-05 18:30:00.000000+0000	HR		HR   raju	20000

(3 rows)

```
cqlsh:employee> begin batch insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (121, 'shyam', 'Software Engineer II', '2022-05-05', 15000.0, 'IT') insert into EmployeeInfo (empid, empname, designation, date_of_joining, salary, dept) values (122, 'babu', 'Software Engineer I', '2022-05-06', 10000.0, 'IT') apply batch;
```

```
cqlsh:employee> select * from EmployeeInfo
```

... ;

empid	date_of_joining	dept	designation	empname	salary
1	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	ramesh	15000
122	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	babu	10000
2	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	suresh	10000
121	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	shyam	15000
3	2022-06-05 18:30:00.000000+0000	HR		HR   raju	20000

(5 rows)

```
cqlsh:employee> update EmployeeInfo set empname
```

```
... ="babu rao" where empid=121;
```

**SyntaxException: line 2:12 no viable alternative at input 'where' (update EmployeeInfo set empname=["babu rao"] where...)**

```
cqlsh:employee> update EmployeeInfo set empname ="babu rao" where empid=121;
```

**SyntaxException: line 1:44 no viable alternative at input 'where' (update EmployeeInfo set empname =["babu rao"] where...)**

```
cqlsh:employee> update EmployeeInfo set empname ="babu rao" where empid=121;
```

```
cqlsh:employee> select * from EmployeeInfo;
```

empid	date_of_joining	dept	designation	empname	salary
1	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	ramesh	15000
122	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	babu	10000
2	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	suresh	10000
121	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	babu rao	15000
3	2022-06-05 18:30:00.000000+0000	HR		HR   raju	20000

(5 rows)

cqlsh:employee> PAGING OFF

...;

**SyntaxException: line 1:0 no viable alternative at input 'PAGING' ([PAGING]...)**

cqlsh:employee> PAGING OFF;

Disabled Query paging.

cqlsh:employee> select \* from EmployeeInfo where empid in (1,2,3) order by salary allow filtering;

**InvalidRequest: Error from server: code=2200 [Invalid query] message="Order by is currently only supported on the clustered columns of the PRIMARY KEY, got salary"**

cqlsh:employee> create index EmployeeInfo(salary);

**SyntaxException: line 1:25 mismatched input '(' expecting K\_ON (create index EmployeeInfo[(]...)**

cqlsh:employee> create index on EmployeeInfo(salary);

cqlsh:employee> select \* from EmployeeInfo where empid in (1,2,3) order by salary allow filtering;

**InvalidRequest: Error from server: code=2200 [Invalid query] message="Order by is currently only supported on the clustered columns of the PRIMARY KEY, got salary"**

cqlsh:employee> alter table EmployeeInfo add projects set<text>;

cqlsh:employee> select \* from EmployeeInfo;

empid	date_of_joining	dept	designation	empname	projects	salary
1	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	ramesh	null	15000
122	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	babu	null	10000
2	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	suresh	null	10000
121	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	babu rao	null	15000
3	2022-06-05 18:30:00.000000+0000	HR		HR   raju	null	20000

(5 rows)

cqlsh:employee> update EmployeeInfo set projects = { 'weather prediction', 'stock prediction' } where empid=1;

cqlsh:employee> update EmployeeInfo set projects = { 'dashboard app' } where empid=2;

cqlsh:employee> update EmployeeInfo set projects = { 'human behavior stats' } where empid=3;

cqlsh:employee> select \* from EmployeeInfo;

empid	date_of_joining	dept	designation	empname	projects	salary
1	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	ramesh	{ 'stock prediction', 'weather prediction' }	15000
122	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	babu		null   10000
2	2022-05-05 18:30:00.000000+0000	IT	Software Engineer I	suresh	{ 'dashboard app' }	10000
121	2022-05-04 18:30:00.000000+0000	IT	Software Engineer II	babu rao		null   15000
3	2022-06-05 18:30:00.000000+0000	HR		HR   raju	{ 'human behavior stats' }	20000

(5 rows)

cqlsh:employee>